



分布式系统

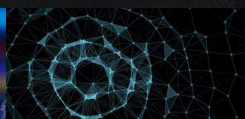
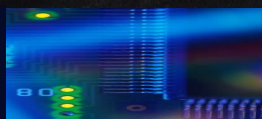
Distributed Systems

陈鹏飞
计算机学院

chenpf7@mail.sysu.edu.cn

办公室：学院楼310

主页：<http://sdcs.sysu.edu.cn/node/3747>



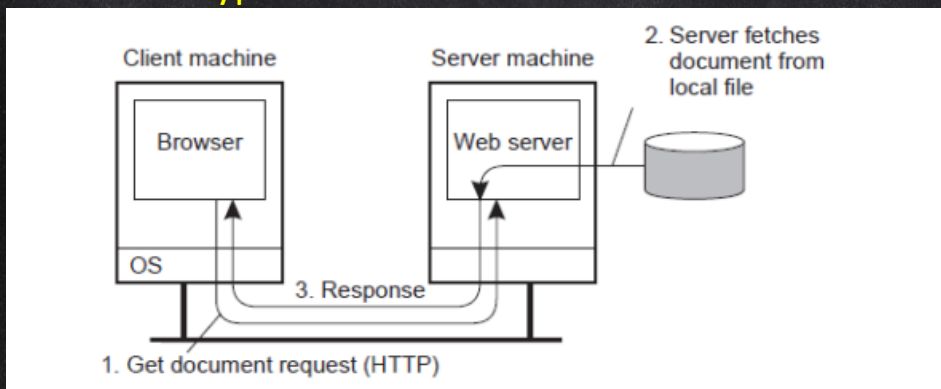
第十二讲 — 分布式Web系统



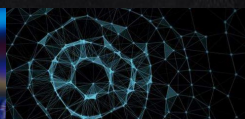
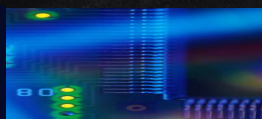
Distributed Web-Based Systems

➤ Essence

The WWW is a huge client-server system with millions of servers; each server hosting thousands of hyperlinked documents:



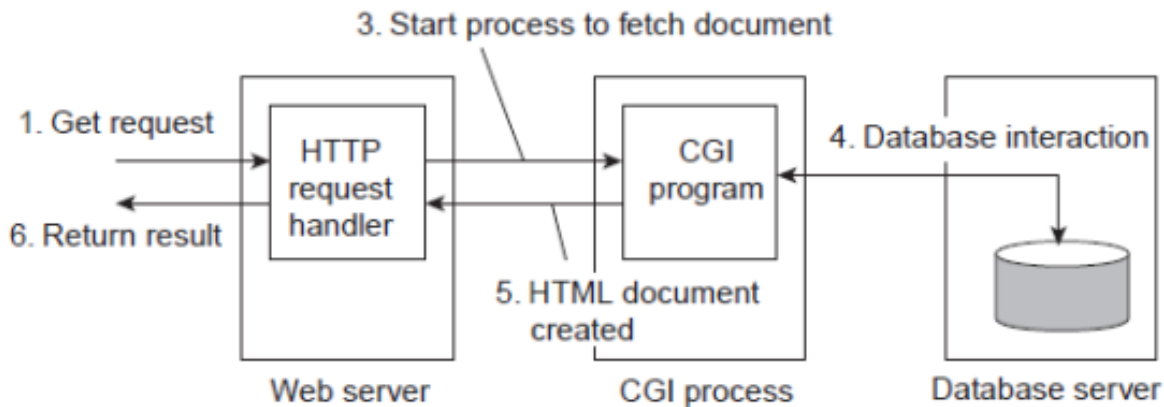
- Documents are generally represented in text (plain text, HTML, XML)
- Alternative types: images, audio, video, but also applications (PDF, PS)
- Documents contain scripts that are executed by the client-side software



Multi-tiered Architectures

➤ Observation

Conventionally, web sites were organized into three tiers.

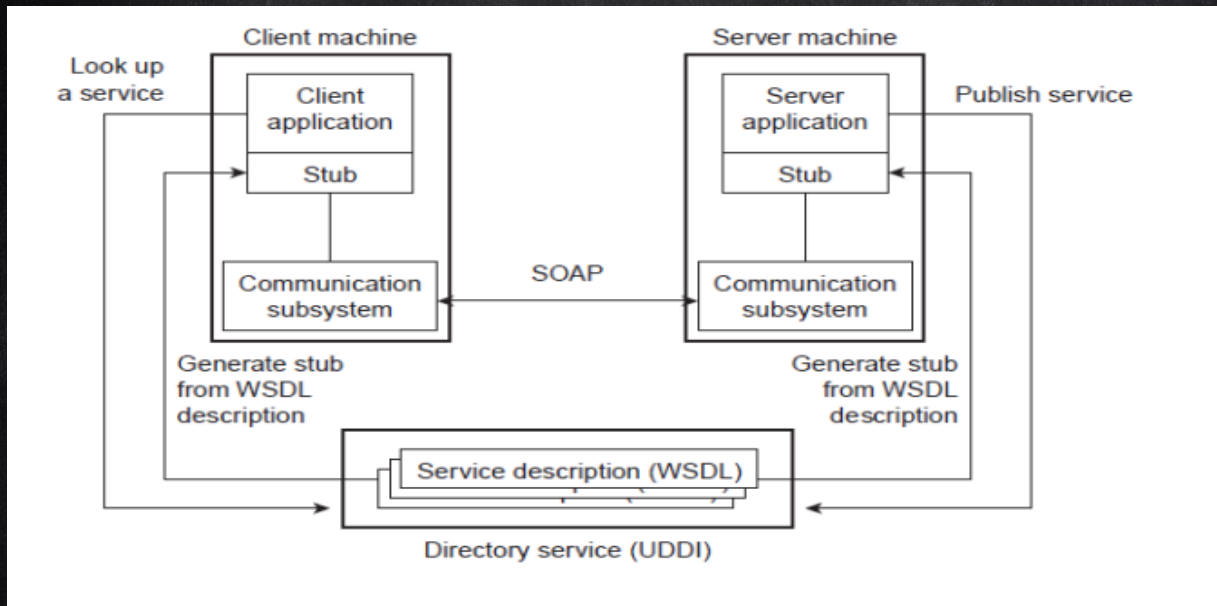




Web services

➤ Observation

At a certain point, people started recognizing that it is was more than just **user** <-> **site** interaction: sites could offer services to other sites => standardization is then badly needed.

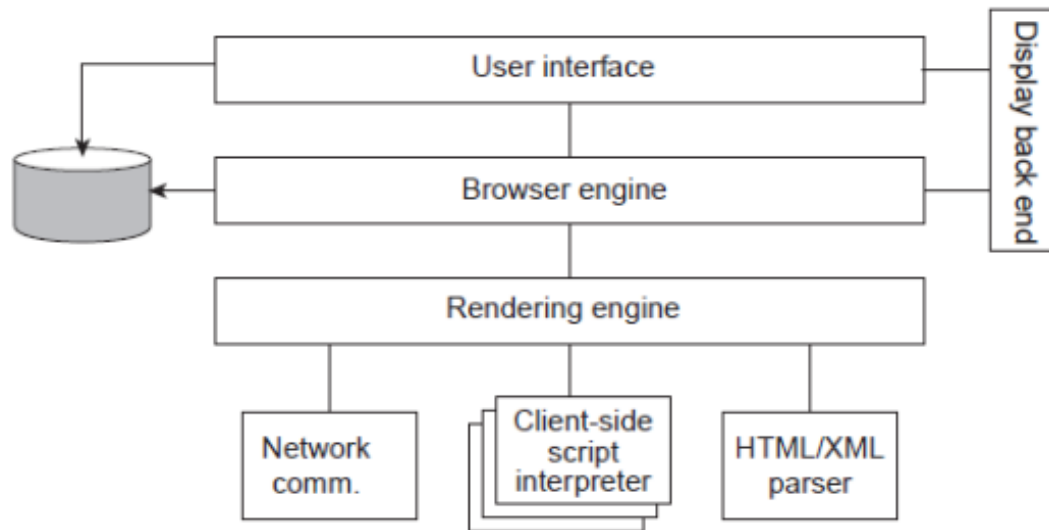




Clients: Web browsers

➤ Observation

browsers form the Web's most important client-side software. They are used to be simple, but that is long ago.

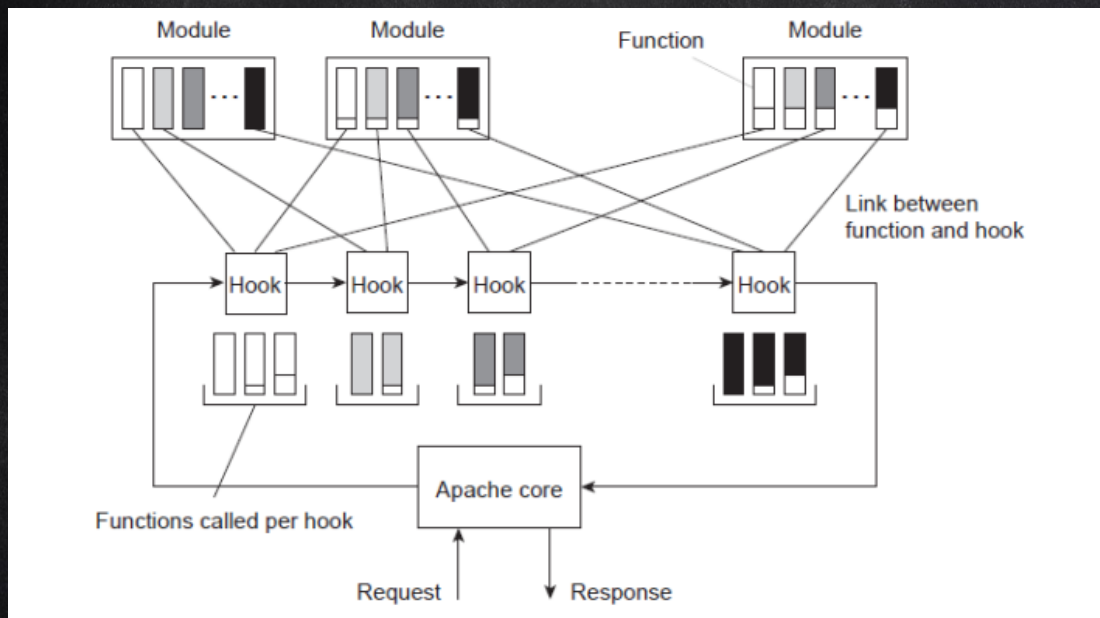




Apache Web Server

➤ Observation

More than 70% of all Web sites are based on Apache. The server is internally organized more or less according to the steps needed to process an HTTP request:

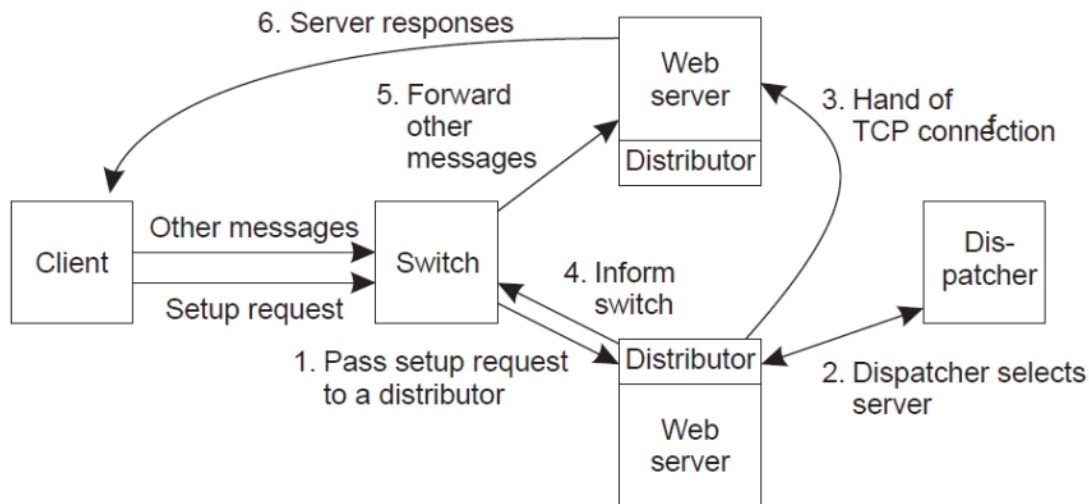




Server Clusters (2/2)

➤ Question

Why can content-aware distribution be so much better?





Communication (1/2)

➤ Essence

Communication in the Web is generally based on HTTP; a relatively simple client-server transfer protocol having the following request messages:

Operation	Description
Head	Request to return the header of a document
Get	Request to return a document to the client
Put	Request to store a document
Post	Provide data that are to be added to a document (collection)
Delete	Request to delete a document



Communication (2/2)

Header	C/S	Contents
Accept	C	The type of documents the client can handle
Accept-Charset	C	The character sets are acceptable for the client
Accept-Encoding	C	The document encodings the client can handle
Accept-Language	C	The natural language the client can handle
Authorization	C	A list of the client's credentials
WWW-Authenticate	S	Security challenge the client should respond to
Date	C+S	Date and time the message was sent
ETag	S	The tags associated with the returned document
Expires	S	The time for how long the response remains valid
From	C	The client's e-mail address
Host	C	The TCP address of the document's server
If-Match	C	The tags the document should have
If-None-Match	C	The tags the document should not have
If-Modified-Since	C	Tells the server to return a document only if it has been modified since the specified time
If-Unmodified-Since	C	Tells the server to return a document only if it has not been modified since the specified time
Last-Modified	S	The time the returned document was last modified
Location	S	A document reference to which the client should redirect its request
Referer	C	Refers to client's most recently requested document
Upgrade	C+S	The application protocol sender wants to switch to
Warning	C+S	Information about status of the data in the message



SOAP

Simple Object Access Protocol: Based on XML, this is the standard protocol for communication between Web services.

- SOAP is bound to an underlying protocol (i.e., it is not independent from its carrier, HTTP、SMTP)
- **Conversational exchange style:** Send a document one way, get a filled-in response back.
- **RPC-style exchange:** Used to invoke a Web service.



A Note on XML

➤ Observation

XML has the advantage of allowing **self describing documents**.

```
env:Envelope xmlns:env="http://www.w3.org/2003/05/soap-envelope">
  <env:Header>
    <n:alertcontrol xmlns:n="http://example.org/alertcontrol">
      <n:priority>1</n:priority>
      <n:expires>2001-06-22T14:00:00-05:00</n:expires>
    </n:alertcontrol>
  </env:Header>
  <env:Body>
    <m:alert xmlns:m="http://example.org/alert">
      <m:msg>Pick up Mary at school at 2pm</m:msg>
    </m:alert>
  </env:Body>
</env:Envelope>
```




Naming: URL

➤ URL

Uniform Resource Locator tells how and where to access a resource.

Scheme	Host name	Pathname
http	:// www.cs.vu.nl	/home/steen/mbox

(a)

Scheme	Host name	Port	Pathname
http	:// www.cs.vu.nl	: 80	/home/steen/mbox

(b)

Scheme	Host name	Port	Pathname
http	:// 130.37.24.11	: 80	/home/steen/mbox

(c)

http	HTTP	http://www.cs.vu.nl:80/globe
mailto	Mail	mailto:steen@cs.vu.nl
ftp	FTP	ftp://ftp.cs.vu.nl/pub/minix/README
file	Local file	file:/edu/book/work/chp/11/11
data	Inline data	data:text/plain;charset=iso-8859-7, %e1%e2%e3
telnet	Remote login	telnet://flits.cs.vu.nl



Web Proxy Caching

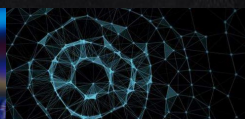
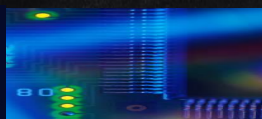
➤ Basic Idea

Sites install a separate proxy server that handles all outgoing requests. Proxies subsequently cache incoming documents. Cache-consistency protocols:

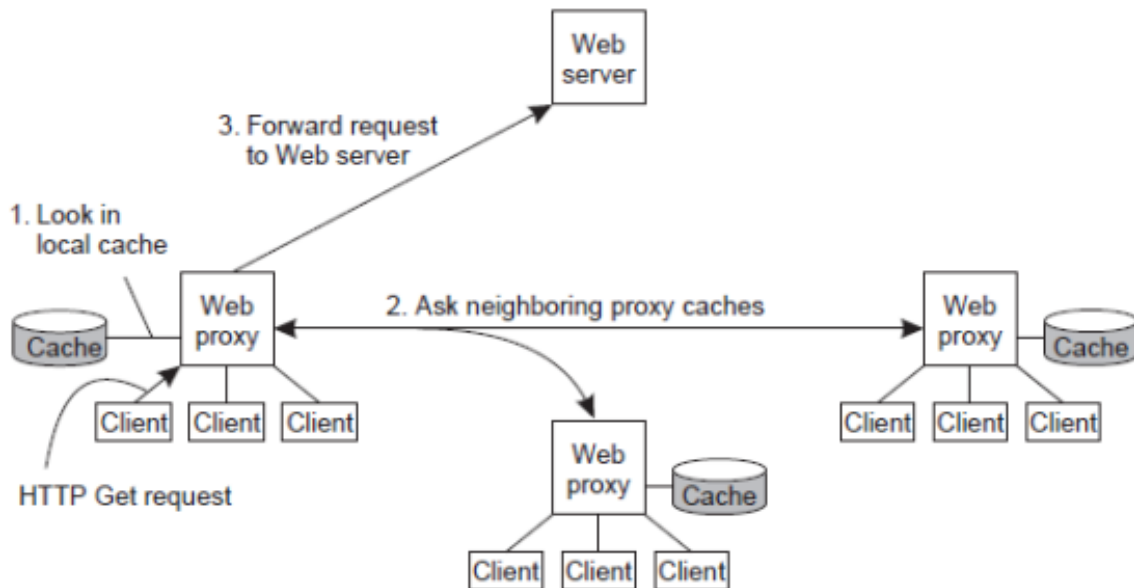
- Always verify validity by contacting server
- Age-based consistency:

$$T_{expire} = \alpha \cdot (T_{cached} - T_{last_modified}) + T_{cached}$$

- Cooperative caching, by which you first check your neighbors on a cache miss:



Web Proxy Caching

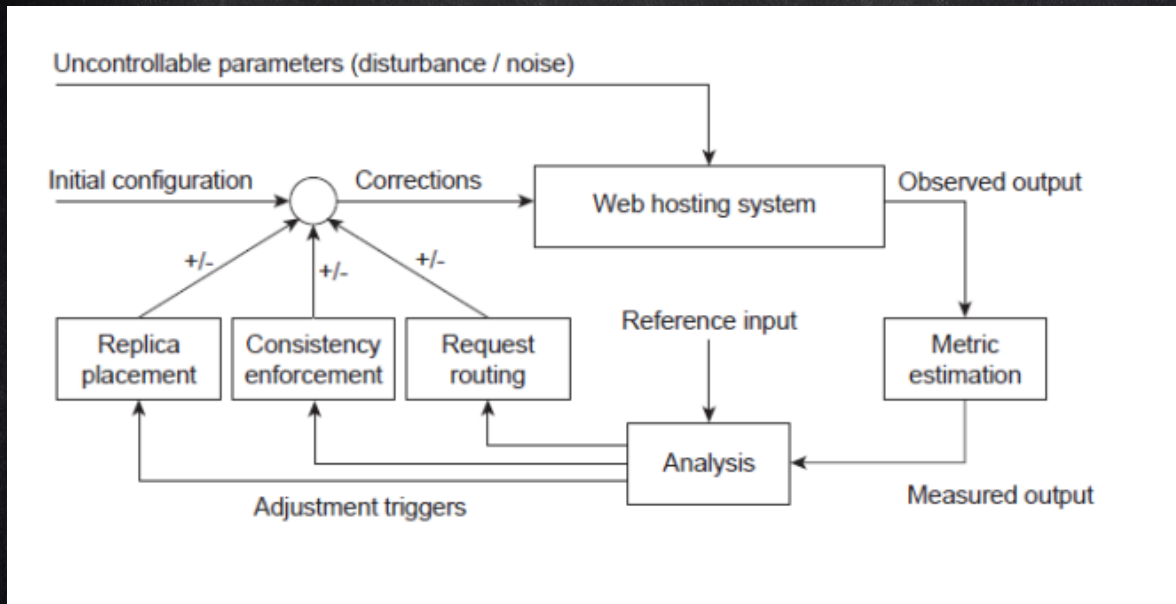




Replication in Web Hosting Systems

➤ Observation

Sites install a separate proxy server that handles all outgoing requests. Proxies subsequently cache incoming documents. Cache-consistency protocols:

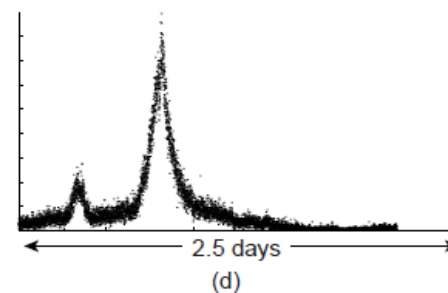
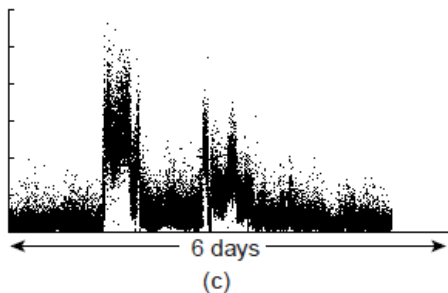
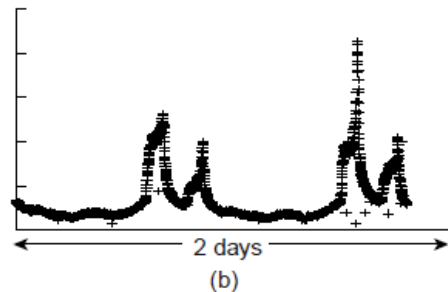
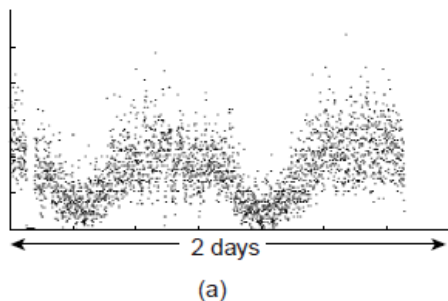




Handling Flash Crowds

➤ Observation

We need dynamic adjustment to balance resource usage. Flash crowds introduce a serious problem:

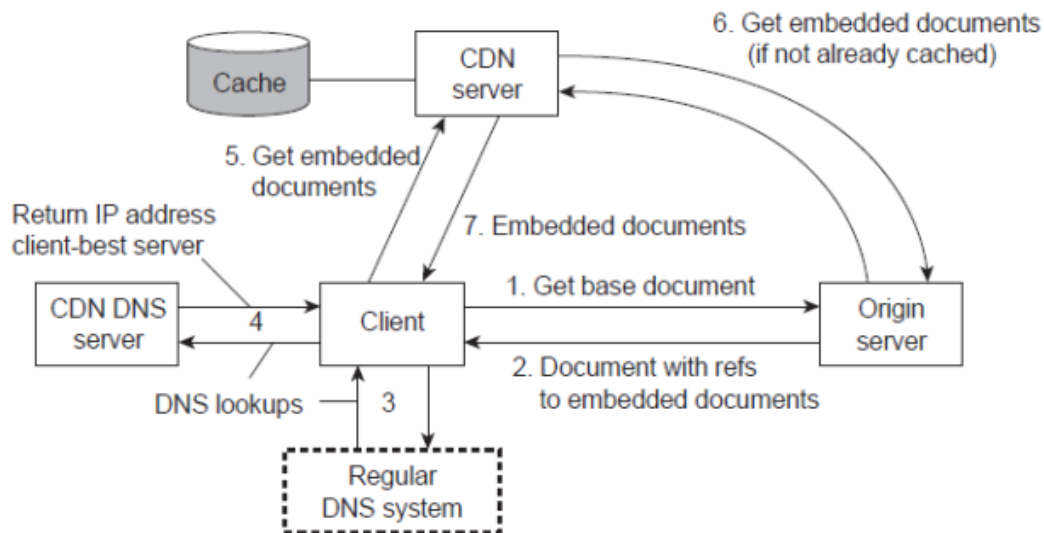




Server Replication

➤ Content Delivery Network

CDNs act as Web hosting services to replicate documents across the Internet providing their customers guarantees on high availability and performance (example: Akamai).

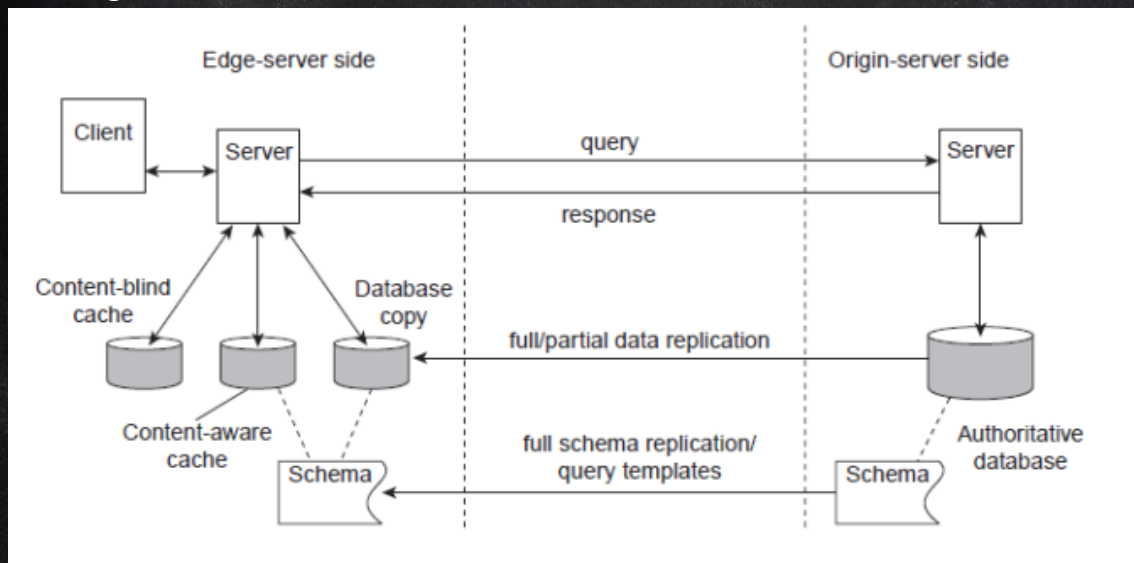




Replication of Web Apps. (1/3)

➤ Observation

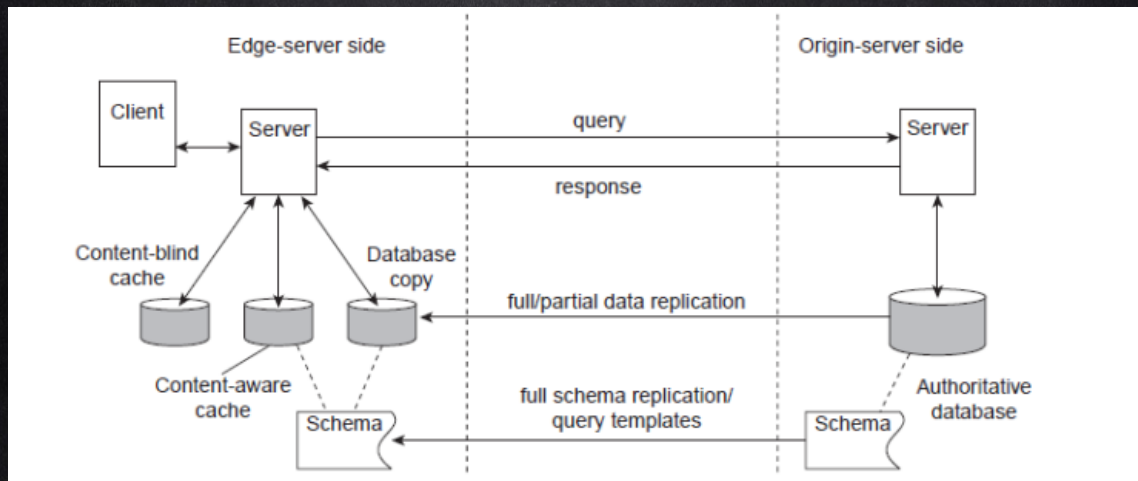
Replication becomes more difficult when dealing with databases and such. No single best solution.



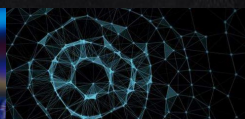
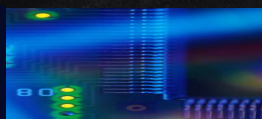
Assumption: Updates are carried out at origin server, and propagated to edge servers.



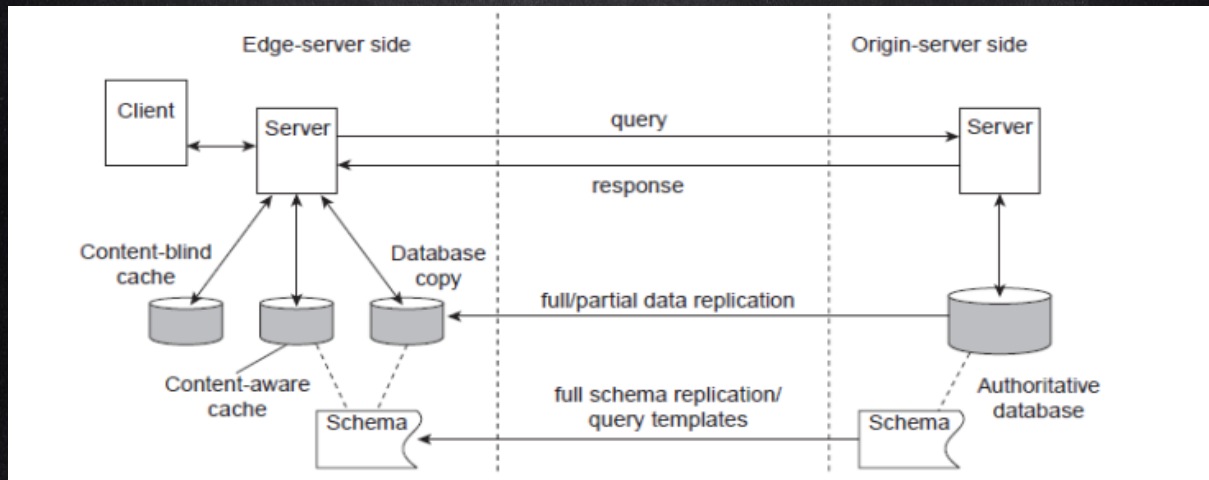
Replication of Web Apps. (2/3)



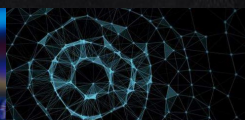
- **Full replication:** high read/write ratio, often in combination with complex queries. **Note:** replication may possibly speed-down performance when R/W ratio goes down.
- **Partial replication:** high read/write ratio, but in combination with simple queries



Replication of Web Apps. (3/3)



- **Content-aware caching:** Check for queries at local database, and subscribe for invalidations at the server. Works good with range queries and complex queries.
- **Content-blind caching:** Simply cache the result of previous queries. Works great with simple queries that address unique results (e.g., no range queries).



谢谢!